

Chahat Raj

PhD in Computer Science | George Mason University

🌐 chahatraj.github.io @ craj@gmu.edu 🎓 Google Scholar 📄 github.com/chahatraj 🐦 [chahatsaidit](https://twitter.com/chahatsaidit)

Research Interests

My research focuses on NLP, particularly the evaluation and alignment of large language models. I study socio-cultural biases and other harms in multilingual and multimodal systems and develop methods to mitigate them, advancing trust-worthy and responsible AI. My work has been published at ACL, EMNLP, AIES, AAAI, and other venues.

Education

May 2027	George Mason University	Fairfax, Virginia
Aug 2022	Ph.D. in Computer Science (GPA: 3.83) Advisors: Ziwei Zhu, Antonios Anastasopoulos	
Aug 2021	Delhi Technological University	Delhi, India
Aug 2019	Masters in Information Systems (Thesis) Advisor: Priyanka Meel	
May 2019	Indira Gandhi Delhi Technical University for Women	Delhi, India
Aug 2015	Bachelors in Computer Science & Engineering	

Experience

Aug 2026	Amazon	Seattle, WA
May 2026	Applied Scientist Intern Mentor: Ren Pang , Manager: Zheng Qi > A parameter-efficient safety steering technique for Bedrock foundation models.	
Aug 2025	George Mason University, University of Washington	Fairfax, VA
May 2024	Graduate Research Assistant Advisors: Ziwei Zhu , Antonios Anastasopoulos , Aylin Caliskan > Societal simulations, model unlearning, and cognitive alignment for bias mitigation.	
Sep 2024	University of Washington Tech Policy Lab, The Information School	Seattle, WA
May 2024	Visiting Scholar Advisor: Aylin Caliskan > Bias, cognition, and mitigation in Vision Language Models.	
Aug 2023	George Mason University School of Computer Science	Fairfax, VA
May 2023	Graduate Research Assistant Advisor: Ziwei Zhu > Fairness, interpretability, and robustness in transformer-based fake news detection.	
Jul 2022	University of Technology Sydney School of Computer Science	Remote
Jun 2021	Visiting Scholar Advisor: Mukesh Prasad > Cyberbullying, AI for Aborigines, modeling disaster fundraising, and ML for clinical diagnosis.	

Selected Publications

[full list here](#)

- [C.13] **Debias It Yourself: Teaching LLMs Cognitive Bias-Mitigation Interventions** [PDF | Code]
[Chahat Raj](#), Sina Mansouri, Aylin Caliskan, Antonios Anastasopoulos, Ziwei Zhu
Under Review [In Submission]
- [C.12] **VIGNETTE: Socially Grounded Bias Evaluation for Vision-Language Models** [PDF | Code]
[Chahat Raj](#), Bowen Wei, Aylin Caliskan, Antonios Anastasopoulos, Ziwei Zhu
In Proceedings of the Association for Computational Linguistics: ACL 2026, San Diego, California, USA [ACL Main'26]
- [C.11] **Talent or Luck? Evaluating Attribution Bias in Large Language Models** [PDF | Code]
[Chahat Raj](#), Mahika Banerjee, Jinhao Pan, Aylin Caliskan, Antonios Anastasopoulos, Ziwei Zhu
In Findings of the Association for Computational Linguistics: ACL 2026, San Diego, California, USA [ACL Findings'26]
- [C.10] **Knowing Bias, Doing Better: Mitigating Social Bias in LLMs via Know-Bias Neuron Enhancement** [PDF | Code]
Jinhao, [Chahat Raj](#), Anjishnu Mukherjee, Sina Mansouri, Bowen Wei, Shloka Yada, Ziwei Zhu
International Conference in Machine Learning: ICML 2026, Seoul, South Korea [ICML'26]
- [C.9] **Measuring South Asian Biases in Large Language Models** [PDF | Code]
Mamnuya Rinki, [Chahat Raj](#), Anjishnu Mukherjee, Ziwei Zhu
TrustNLP at ACL 2026 [TrustNLP'26]
- [C.8] **Discovering Bias Associations through Open-Ended LLM Generations** [PDF | Code]
Jinhao Pan, [Chahat Raj](#), Ziwei Zhu
In Proceedings of the AAAI Conference on Artificial Intelligence [AAAI'26]

- [C.7] **What's Not Said Still Hurts: A Description-Based Evaluation Framework for Measuring Social Bias in LLMs**
Jinhao Pan, [Chahat Raj](#), Ziyu Yao, Ziwei Zhu [PDF | Code]
In Findings of the Association for Computational Linguistics: EMNLP 2025, Suzhou, China [EMNLP Findings'25]
- [C.6] **Toward Inclusive Language Models: Sparsity-Driven Calibration for Systematic and Interpretable Mitigation of Social Biases in LLMs** [PDF | Code]
Prommy Sultana Hossain, [Chahat Raj](#), Ziwei Zhu, Jessica Lin, Emanuela Marasco
In Findings of the Association for Computational Linguistics: EMNLP 2025, Suzhou, China [EMNLP Findings'25]
- [C.5] **BiasDora: Exploring Hidden Biased Associations in Vision-Language Models** [PDF | Code]
[Chahat Raj](#), Anjishnu Mukherjee, Aylin Caliskan, Antonios Anastasopoulos, Ziwei Zhu
In Findings of the Association for Computational Linguistics: EMNLP 2024, Miami, Florida, USA [EMNLP Findings'24]
- [C.4] **Breaking Bias, Building Bridges: Evaluation and Mitigation of Social Biases in LLMs via Contact Hypothesis**
[Chahat Raj](#), Anjishnu Mukherjee, Aylin Caliskan, Antonios Anastasopoulos, Ziwei Zhu [PDF | Code]
Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society [AIES'24]
- [C.3] **Global Voices, Local Biases: Socio-cultural Prejudices across Languages** [PDF | Code]
[Chahat Raj](#)*, Anjishnu Mukherjee*, Ziwei Zhu, Antonios Anastasopoulos
In Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, Singapore [EMNLP Main'23]
- [C.2] **SALSA: Salience-Based Switching Attack for Adversarial Perturbations in Fake News Detection Models**
[Chahat Raj](#)*, Anjishnu Mukherjee*, Hemant Purohit, Antonios Anastasopoulos, Ziwei Zhu [PDF | Code]
European Conference on Information Retrieval [ECIR'24]
- [C.1] **True and Fair: Robust and Unbiased Fake News Detection via Interpretable Machine Learning** [PDF | Code]
[Chahat Raj](#), Anjishnu Mukherjee, Ziwei Zhu
Proceedings of the AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society [AIES'23]

Selected Projects

- | | | |
|---|--|--------------------------|
| > AI, Power & Society | Auditing Diffusion Models for Generational and Representational Errors | [Report] |
| > Computer Vision | Cultural Image Translation for Human Figures | [Report] |
| > AI: Ethics, Policy & Society | Implicit Association Test in Multimodal Generative Models | [Report] |

Teaching

- | | | |
|--|--|--------------------|
| > CS 484 Theory/Application Data Mining, GMU | <i>Inst.: Keren Zhou</i> | Spring'26 |
| > CS 584 Theory/Application Data Mining, GMU | <i>Inst.: Ziwei Zhu</i> | Fall'25 |
| > CS 108 Introduction to Computer Programming, GMU | <i>Inst.: Sara Lynn Gopalakrishna, Bella Yin</i> | Spring'24 |
| > CS 478 Natural Language Processing, GMU | <i>Inst.: Antonios Anastasopoulos</i> | Fall'23 |
| > COMP 502 Mathematical Foundations of Computing, GMU | <i>Inst.: William D Ellis</i> | Fall'23 |
| > CS 112 Introduction to Python Programming, GMU | <i>Inst.: Shahnaz Kamberi</i> | Fall'22, Spring'23 |

Student Mentorship

sub-mentoring with Ziwei, Antonis, Aylin

- | | |
|---|--|
| > Daksha Rajagopal [TJHSST → Purdue CS] | Modeling normative ethics and cross-cultural moral reasoning in VLMs |
| > Mahika Banerjee [TJHSST → KU Leuven] | Measuring attribution bias, coauthor on the ACL Findings 2026 paper |
| > Mamnuya Rinki [MS Thesis, GMU → IntraFi] | Mitigating South Asian Biases, coauthor on TrustNLP @ ACL 2026 paper |
| > Angela Liu [TJHSST] | Assessing moral attributional biases in ambiguous contexts in LLMs |
| > Yongxu Sun [MS, UW → PhD, UW] | Inoculation Prompting by in-context learning to debias LLMs' decisions |
| > Diwita Banerjee [MS, GMU → DoorDash] | Assessing cuisine biases; scientific figure explanation issues in VLMs |
| > Rinki, Sharanya, Aksh [Advanced NLP] | Intersectional biases in LLMs across Indo-Aryan languages |

Awards

- | | |
|--|--|
| > SAC Highlight Award @ ACL, 2026 | Vignette selected as top 15% of accepted papers. |
| > Outstanding Academic Award @ GMU, 2026 | Received an Outstanding Academic Achievement Award. |
| > Graduate Teaching Award @ GMU, 2026 | Received an Outstanding Graduate Teaching Assistant Award. |
| > Outstanding Reviewer Award @ EACL 2026 | Received an Outstanding Reviewer Award at EACL 2026. |
| > OpenAI Researcher Access Program Award 2025 | Funded \$1000 worth of OpenAI API credits. |
| > Travel Grant by AAAI/ACM AIES 2024 & 2023 | Funded \$2000 to attend AIES in San Jose and Montreal. |
| > CAHMP Fellowship @ GMU, 2024 | Summer research on social biases in question answering. |
| > Best Paper Award @ MASC-SLL 2024 | A Psychological View to Social Bias in LLMs. |
| > Summer Graduate Research Award @ GMU, 2023 | Explainability and Robustness of Debiasing. |
| > Research Excellence Award @ DTU 2023 & 2022 | Received three Commendable Research Awards with INR 150k. |
| > Graduate Scholarship @ DTU 2019 | Received INR 297k by AICTE for qualifying GATE. |

Talks

> VIGNETTE	<i>Talk at ACL</i>	2026 · San Diego, CA
> LLM Ethics: Bias and Mitigation	<i>Invited Lecture, CS 782 Advanced ML</i>	2024 · GMU, VA
> LLM Ethics: Bias and Mitigation	<i>Invited Lecture, CS 678 Advanced NLP</i>	2024 · GMU, VA
> Breaking Bias, Building Bridges	<i>Talk at AAAI/ACM AIES</i>	2024 · San Jose, CA
> Breaking Bias, Building Bridges	<i>Talk at SouthNLP</i>	2024 · Emory, GA
> Breaking Bias, Building Bridges	<i>Talk at MASC-SLL</i>	2024 · Johns Hopkins, MD
> True and Fair	<i>Talk at AAAI/ACM AIES</i>	2023 · Montreal, QC
> Cyberbully Monitoring	<i>ML4AU Research Showcase</i>	2021

Program Committee

- > **Reviewer** NeurIPS'26, ARR'23-26, CHI'26, AAAI'26, IASEAI'26, COLM'25-26, TrustNLP'24-26, LTEDI @ EACL'24, AIES'23-26, CSCW'24, NeurIPS'23, ACM TIST