
Our models are as ethical as we are!

Models do not mirror humans,
yet they can amplify, suppress, or transform our prejudice.

THESIS STATEMENT

Bias has long shaped society, and social psychology has spent decades studying ways to identify and reduce it.

By positioning social psychology as a foundation for bias research in NLP, my dissertation moves beyond narrow benchmark-driven association tests and targeted debiasing approaches to

- (1) **discover** hidden stereotypes,
- (2) **evaluate** how these stereotypes manifest in model behavior,
- (3) **mitigate** these biases with generalizable methods, and
- (4) move beyond isolated model behavior to **interactive systems** toward more responsible and trustworthy NLP.

My Matryoshka Model

1. BIAS DISCOVERY

EMNLP '23 Main

***Global Voices, Local Biases**

EMNLP '24 Findings

***BiasDora**

AAAI '26

BAD-F

Exploration

Evaluation

Mitigation

***Breaking Bias**

AIES '24

***Debias It Yourself**

Submission (EACL)

Towards Inclusive LMs

EMNLP '25 Findings

Knowing Bias, Doing Better

ICML '26

3. BIAS MITIGATION

2. BEHAVIORAL EVALUATION

ACL '26 Main

***Vignette**

ACL '26 Findings

***Talent or Luck**

TrustNLP '26

Purdah and Patriarchy

EMNLP '25 Findings

What's Not Said Still Hurts

*(Paper) = First authored

My Research Landscape

